

# State Farm Project Model 2

July 29, 2019

## 1 State Farm Classification Project: Model 2

```
In [1]: import pandas as pd
import numpy as np
from sklearn.model_selection import cross_validate, train_test_split, GridSearchCV
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
%matplotlib inline
import matplotlib.pyplot as plt
from sklearn import metrics
```

```
In [3]: train = pd.read_csv("exercise_02_train.csv")
```

```
In [4]: train.head()
```

```
Out[4]:
```

	x0	x1	x2	x3	x4	x5	x6	\
0	0.198560	74.425320	67.627745	-3.095111	-6.822327	19.048071	-0.362378	
1	-29.662621	24.320711	-48.205182	1.430339	-6.552206	4.263074	6.551412	
2	15.493759	-66.160459	50.512903	-2.265792	14.428578	2.509323	-6.707536	
3	-19.837651	33.210943	53.405563	1.079462	11.364251	-1.064581	9.308857	
4	11.896655	-26.717872	-17.758176	1.692017	21.553537	-5.852097	-0.857435	

	x7	x8	x9	...	x91	x92	x93	x94	\
0	-10.699174	-22.699791	-1.561262	...	0.800948	1.553846	asia	-1.093926	
1	4.265483	1.245095	2.246814	...	2.031707	7.544422	asia	-3.659541	
2	3.820842	-11.100833	-1.459825	...	-0.992474	1.385799	america	1.299144	
3	9.266076	14.552959	-2.012755	...	-1.157845	6.036804	asia	0.521396	
4	-2.186940	18.075272	-1.404618	...	-3.045511	-1.719337	asia	1.526071	

	x95	x96	x97	x98	x99	y
0	16.202557	26.238591	-2.125570	9.644466	1.237667	0
1	29.674259	-15.141647	-36.030599	5.820376	1.952183	1
2	33.018090	-19.914894	26.212736	2.372690	0.558988	1
3	9.664095	-27.197636	19.221130	13.382712	0.214462	0
4	-25.608326	33.383803	-5.703269	-11.023730	-1.191319	0

[5 rows x 101 columns]

```

In [5]: def unique(list1):

        # insert the list to the set
        list_set = set(list1)
        # convert the set to the list
        unique_list = (list(list_set))
        for x in unique_list:
            print(x)

        unique(train.dtypes)

int64
float64
object

In [6]: predictCols = list(train)
        predictCols.remove('y')

In [7]: for col in predictCols:
        if train[col].dtype in [np.float64, np.int64]:
            train[col].fillna(train[col].mean(skipna = True), inplace=True)

In [8]: # Ensure no remaining na's
        numericCols = train.select_dtypes(include='number').columns
        naVals = train[numericCols].isna().sum().sort_values()
        naVals.sum()

Out[8]: 0

In [9]: objectCols = train.select_dtypes(include='object').columns
        print(objectCols)

Index(['x34', 'x35', 'x41', 'x45', 'x68', 'x93'], dtype='object')

In [10]: train.x34.fillna(train.x34.mode()[0], inplace=True)
         unique(train['x34'])

tesla
bmw
volkswagon
mercedes
Honda
Toyota
ford
chevrolet
chrystler
nissan

```

```
In [11]: train.x35.replace(['thursday', 'thur'], ['thursday', 'thursday'], inplace=True)
train.x35.replace(['wed'], ['wednesday'], inplace=True)
train.x35.replace(['fri'], ['friday'], inplace=True)
train.x35.fillna(train.x35.mode()[0], inplace=True)
unique(train['x35'])
```

```
monday
thursday
tuesday
wednesday
friday
```

```
In [12]: # Convert currency column to float, remove nan's
train['x41'] = train['x41'].astype(str)
train['x41'] = train['x41'].map(lambda x: x.lstrip('$'))
train['x41'] = train['x41'].astype(np.float16)
train['x41'].fillna(0, inplace=True) # probably safer to replace nan's with 0, not me

print(train['x41'].isna().sum())
```

```
0
```

```
In [13]: # Convert percentage column to float, remove nan's
train['x45'] = train['x45'].astype(str)
train['x45'] = train['x45'].map(lambda x: x.rstrip('%'))
train['x45'] = train['x45'].astype(np.float16)
train['x45'].fillna(train['x45'].mean(skipna = True), inplace=True) # since very few

print(train['x41'].isna().sum())
```

```
0
```

```
In [14]: # Month Column
train.x68.replace(['Dev'], ['Dec'], inplace=True) # because I'm OCD
train.x68.replace(['sept.'], ['Sep'], inplace=True)
train.x68.replace(['January'], ['Jan'], inplace=True)
train.x68.replace(['July'], ['Jul'], inplace=True)
train.x68.fillna(train.x68.mode()[0], inplace=True)

unique(train['x68'])
```

```
Jun
Apr
Aug
Feb
Jul
```

Oct  
Dec  
Jan  
Mar  
Sep  
Nov  
May

```
In [15]: # Region
train.x93.replace(['euorpe'], ['europe'], inplace=True)
train = train[pd.isna(train['x93']) == False]
print(train['x93'].isna().sum())
# Region seems significant, and there's only 7 NA's, so remove rows with this as NA
```

0

```
In [16]: # Check if target has na's
print(train['y'].isna().sum())
```

0

```
In [17]: train = pd.get_dummies(train)
```

```
In [18]: # Ensure we converted all non-numeric columns to numeric
train.select_dtypes(include='object').columns
```

```
Out[18]: Index([], dtype='object')
```

```
In [19]: train.describe()
```

```
Out[19]:
```

	x0	x1	x2	x3	x4 \
count	39993.000000	39993.000000	39993.000000	39993.000000	39993.000000
mean	3.447752	-7.788416	1.704644	-0.072832	0.121980
std	16.245334	37.012224	38.382930	1.503022	16.289301
min	-60.113902	-157.341119	-163.339956	-6.276969	-61.632319
25%	-7.595295	-32.731869	-24.141082	-1.087780	-10.896141
50%	3.446322	-7.987507	1.959477	-0.062721	0.105307
75%	14.266326	16.848201	27.511371	0.940330	11.076726
max	75.311659	153.469221	154.051060	5.837559	65.949709

	x5	x6	x7	x8	x9 \
count	39993.000000	39993.000000	39993.000000	39993.000000	39993.000000
mean	-0.607009	0.035852	-0.052430	-2.911144	-0.024524
std	15.583132	9.040667	6.952184	13.148182	2.939696
min	-62.808995	-35.060656	-26.736717	-53.735586	-11.497395
25%	-11.181510	-6.089227	-4.746572	-11.722590	-2.003827

50%	-0.576660	0.044975	-0.037833	-2.940961	-0.054184
75%	9.954957	6.100325	4.636585	5.857648	1.954809
max	63.424046	45.053946	34.267792	66.936936	11.271939

	...	x68_Jul	x68_Jun	x68_Mar	x68_May \
count	...	39993.000000	39993.000000	39993.000000	39993.000000
mean	...	0.277199	0.231516	0.010777	0.119221
std	...	0.447621	0.421806	0.103252	0.324052
min	...	0.000000	0.000000	0.000000	0.000000
25%	...	0.000000	0.000000	0.000000	0.000000
50%	...	0.000000	0.000000	0.000000	0.000000
75%	...	1.000000	0.000000	0.000000	0.000000
max	...	1.000000	1.000000	1.000000	1.000000

	x68_Nov	x68_Oct	x68_Sep	x93_america	x93_asia \
count	39993.000000	39993.000000	39993.000000	39993.000000	39993.000000
mean	0.003776	0.022604	0.087140	0.078289	0.885555
std	0.061331	0.148639	0.282044	0.268629	0.318355
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	1.000000
50%	0.000000	0.000000	0.000000	0.000000	1.000000
75%	0.000000	0.000000	0.000000	0.000000	1.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

	x93_europe
count	39993.000000
mean	0.036156
std	0.186681
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

[8 rows x 127 columns]

## 1.1 Now check class imbalance

```
In [20]: x = train.y.value_counts()
print(x)

print("% oF Training Set with Positives: " + "{0:.0%}".format(x[1] / (x[0] + x[1])))
print("% oF Training Set with Negatives: " + "{0:.0%}".format(x[0] / (x[0] + x[1])))

0    31851
1     8142
Name: y, dtype: int64
% oF Training Set with Positives: 20%
```

% of Training Set with Negatives: 80%

**1.1.1 This class imbalance is not too bad, so we don't need to do resampling...**

## **1.2 Now scale data to normalize**

```
In [21]: from sklearn.preprocessing import StandardScaler

predictCols = list(train)
predictCols.remove('y')

train2 = train.copy()
for col in predictCols:
    scaler = StandardScaler()
    if train2[col].dtype != 'float64': # convert data type to avoid warning
        train2[col] = np.float64(train2[col])
    scaler.fit(train2[[col]])
    train2.loc[:,col] = scaler.transform(train2[[col]])
```

In [ ]:


## **1.3 Now check correlation**

```
In [22]: import seaborn as sns
        %matplotlib inline

# calculate the correlation matrix
corr = train[numericCols].corr()

# plot the heatmap
sns.heatmap(corr,
            xticklabels=corr.columns,
            yticklabels=corr.columns)
```

Out [22]: <matplotlib.axes.\_subplots.AxesSubplot at 0x2194c03b3c8>



```
In [23]: # Attempt to find columns with high-correlation, and remove if necessary
for col in numericCols:
    #print("\n\n" + col)
    q = corr[col].sort_values(ascending = False)
    q = q.drop(col)
    q = q[abs(q) >= 0.5]
    perfectNames = q.index.values
    if len(q) > 0:
        print(col)
```

**1.4** Clearly, no columns in the data set are highly correlated, so no need to remove.

```
In [24]: y = train2['y']
         train2 = train2.drop(['y'], axis=1)
```

```
In [ ]:
```

## 2 Now can use Support Vector Machine model

```
In [25]: from sklearn import svm
```

```
In [92]: # Train test split
         X_train, X_test, y_train, y_test = train_test_split(train2, y, test_size=.2, random_s
```

## 2.1 First try linear kernel, although we don't expect this problem to be linear

```
In [93]: svcclassifier = svm.SVC(kernel='linear', C = 1.0, max_iter = 10000)
        svcclassifier.fit(X_train, y_train)
```

```
C:\Python 3.7\lib\site-packages\sklearn\svm\base.py:244: ConvergenceWarning: Solver terminated
% self.max_iter, ConvergenceWarning)
```

```
Out [93]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
            decision_function_shape='ovr', degree=3, gamma='auto_deprecated',
            kernel='linear', max_iter=10000, probability=False, random_state=None,
            shrinking=True, tol=0.001, verbose=False)
```

```
In [94]: pred_train = svcclassifier.predict(X_train)
        pred_test = svcclassifier.predict(X_test)
```

```
In [95]: # evaluate predictions
        train_accuracy = accuracy_score(y_train, pred_train)
        print("Train Accuracy: %.2f%%" % (train_accuracy * 100.0))

        test_accuracy = accuracy_score(y_test, pred_test.round())
        print("Test Accuracy: %.2f%%" % (test_accuracy * 100.0))
```

Train Accuracy: 71.74%

Test Accuracy: 71.32%


```
In [96]: ## Accuracy is not very good, but at least we're not over-fitting
```

```
In [97]: cm = confusion_matrix(y_test, pred_test)
```

```
sns.set(font_scale=1.4)#for label size
sns.heatmap(cm, annot=True,fmt='g',annot_kws={"size": 16})# font size
```

```
Out [97]: <matplotlib.axes._subplots.AxesSubplot at 0x19431df2908>
```





**2.1.1 Slightly more false negatives than false positives...probably due to the class imbalance**  
 0's to 1's in the training set heavily weighted toward 0's (80%)

## 2.2 Now try the Radial Basis Function kernel

```
In [98]: svcclassifier = svm.SVC(kernel='rbf', C = 1.0, max_iter = 10000, gamma = 'auto')
         svcclassifier.fit(X_train, y_train)

         pred_train = svcclassifier.predict(X_train)
         pred_test = svcclassifier.predict(X_test)

         test_accuracy = accuracy_score(y_test, pred_test)
         print("Test Accuracy: %.2f%%" % (test_accuracy * 100.0))
```

C:\Python 3.7\lib\site-packages\sklearn\svm\base.py:244: ConvergenceWarning: Solver terminated  
 % self.max\_iter, ConvergenceWarning)

Test Accuracy: 98.35%

**2.2.1 Massive improvement by switching to rbf kernel. Increased max\_iter beyond 10,000 didn't have much effect on accuracy, but was much slower**

## **2.3 Now use a grid search with 5-fold cross-validation to find the hyperparameters C and gamma**

In [39]: *# Tune the two main hyperparameters*

```
from sklearn.model_selection import GridSearchCV
```

```
def svc_param_selection(X, y, nfolds):
```

```
    Cs = [0.001, 0.01, 0.1, 1, 10]
```

```
    gammas = [0.001, 0.01, 0.1, 1]
```

```
    param_grid = {'C': Cs, 'gamma' : gammas}
```

```
    grid_search = GridSearchCV(svm.SVC(kernel='rbf',max_iter = 1000,gamma = 'auto'),
```

```
    grid_search.fit(X, y)
```

```
    grid_search.best_params_
```

```
    return grid_search.best_params_
```

```
svc_param_selection(X_train, y_train, 5)
```

```
C:\Python 3.7\lib\site-packages\sklearn\svm\base.py:244: ConvergenceWarning: Solver terminated  
% self.max_iter, ConvergenceWarning)
```

```
C:\Python 3.7\lib\site-packages\sklearn\svm\base.py:244: ConvergenceWarning: Solver terminated  
% self.max_iter, ConvergenceWarning)
```

```
C:\Python 3.7\lib\site-packages\sklearn\svm\base.py:244: ConvergenceWarning: Solver terminated  
% self.max_iter, ConvergenceWarning)
```

```
C:\Python 3.7\lib\site-packages\sklearn\svm\base.py:244: ConvergenceWarning: Solver terminated  
% self.max_iter, ConvergenceWarning)
```

```
C:\Python 3.7\lib\site-packages\sklearn\svm\base.py:244: ConvergenceWarning: Solver terminated  
% self.max_iter, ConvergenceWarning)
```

```
C:\Python 3.7\lib\site-packages\sklearn\svm\base.py:244: ConvergenceWarning: Solver terminated  
% self.max_iter, ConvergenceWarning)
```

```
C:\Python 3.7\lib\site-packages\sklearn\svm\base.py:244: ConvergenceWarning: Solver terminated  
% self.max_iter, ConvergenceWarning)
```

```
C:\Python 3.7\lib\site-packages\sklearn\svm\base.py:244: ConvergenceWarning: Solver terminated  
% self.max_iter, ConvergenceWarning)
```

```
C:\Python 3.7\lib\site-packages\sklearn\svm\base.py:244: ConvergenceWarning: Solver terminated  
% self.max_iter, ConvergenceWarning)
```

```
C:\Python 3.7\lib\site-packages\sklearn\svm\base.py:244: ConvergenceWarning: Solver terminated  
% self.max_iter, ConvergenceWarning)
```

```
C:\Python 3.7\lib\site-packages\sklearn\svm\base.py:244: ConvergenceWarning: Solver terminated  
% self.max_iter, ConvergenceWarning)
```

```
C:\Python 3.7\lib\site-packages\sklearn\svm\base.py:244: ConvergenceWarning: Solver terminated  
% self.max_iter, ConvergenceWarning)
```

```
C:\Python 3.7\lib\site-packages\sklearn\svm\base.py:244: ConvergenceWarning: Solver terminated  
% self.max_iter, ConvergenceWarning)
```

```
C:\Python 3.7\lib\site-packages\sklearn\svm\base.py:244: ConvergenceWarning: Solver terminated  
% self.max_iter, ConvergenceWarning)
```

```
C:\Python 3.7\lib\site-packages\sklearn\svm\base.py:244: ConvergenceWarning: Solver terminated
```

[illegible]

[illegible]

[illegible]



```

In [28]: predictCols = list(test)
         for col in predictCols:
             if test[col].dtype in [np.float64, np.int64]:
                 #print(col)
                 test[col].fillna(test[col].mean(skipna = True), inplace=True)

In [29]: test.x34.fillna(test.x34.mode()[0], inplace=True)

In [30]: # Day Column
         test.x35.replace(['thursday', 'thur'], ['thursday', 'thursday'], inplace=True)
         test.x35.replace(['wed'], ['wednesday'], inplace=True)
         test.x35.replace(['fri'], ['friday'], inplace=True)
         test.x35.fillna(test.x35.mode()[0], inplace=True)

In [31]: # Convert currency column to float, remove nan's
         test['x41'] = test['x41'].astype(str)
         test['x41'] = test['x41'].map(lambda x: x.lstrip('$'))
         test['x41'] = test['x41'].astype(np.float16)
         test['x41'].fillna(0, inplace=True) # probably safer to replace nan's with 0, not mean

In [32]: # Convert percentage column to float, remove nan's
         test['x45'] = test['x45'].astype(str)
         test['x45'] = test['x45'].map(lambda x: x.rstrip('%'))
         test['x45'] = test['x45'].astype(np.float16)
         test['x45'].fillna(train['x45'].mean(skipna = True), inplace=True) # since very few u

In [33]: # Month Column
         test.x68.replace(['Dev'], ['Dec'], inplace=True) # because I'm OCD
         test.x68.replace(['sept.'], ['Sep'], inplace=True)
         test.x68.replace(['January'], ['Jan'], inplace=True)
         test.x68.replace(['July'], ['Jul'], inplace=True)
         test.x68.fillna(test.x68.mode()[0], inplace=True)

In [34]: # Region
         test.x93.replace(['euorpe'], ['europe'], inplace=True)
         test.x93.fillna(test.x93.mode()[0], inplace=True)
         print(test['x93'].isna().sum())
         # Region seems significant, and there's only 7 NA's, so remove rows with this as NA

0

In [35]: test = pd.get_dummies(test)

In [36]: # Ensure all columns in test are also in train after the one-hot encoding
         any(elem in list(test) for elem in list(train))

Out[36]: True

```

```
In [38]: # Scale the test data
        for col in test.columns:
            scaler = StandardScaler()
            if test[col].dtype != 'float64': # convert data type to avoid warning
                test[col] = np.float64(test[col])
            scaler.fit(test[[col]])
            test.loc[:,col] = scaler.transform(test[[col]])
```

## 2.5 Now retrain the model with the final hyperparameters using the full training set

```
In [40]: svcclassifier = svm.SVC(kernel='rbf', C = 1.0, gamma = 'auto', probability = True)
        svcclassifier.fit(train2, y)
```

```
Out[40]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
            decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
            max_iter=-1, probability=True, random_state=None, shrinking=True,
            tol=0.001, verbose=False)
```

## 2.6 Now generate the final test output

```
In [41]: final_y = svcclassifier.predict_proba(test) # return class probabilities
```

```
In [43]: final_y = final_y[:,1] # return only the probability of the 1's class
```

```
In [44]: np.savetxt("results2.csv", final_y, delimiter = ",")
```